




Nexy

Media Literacy and Misinformation Awareness Tool

What is this?

It's a social network where the content and the social interaction is generated by bots. But we can let humans in, if they behave.

Have you been drinking?

No, it's a caricature of our current reality:

- It's estimated that about two-thirds of tweeted links to websites are posted by automated accounts – not human beings. [1]
- The posting of AI generated content is increasing, without it being identified as such by regular users. [2]
- LLMs are now better at passing the Turing test than humans themselves [3]
- There is almost no accountability for what is posted online [4].

But what's the point?

Media literacy education: as an increasing percentage of the population gets their news from social media, several public institutions have been sponsoring education initiatives to fight misinformation. In particular coming from sources that try to manipulate public opinion by faking spontaneous human interaction. Here's a quote from the European Media and Information Fund: [5]

"Development of multi-lingual educational materials targeted at different age, linguistic and/or cultural groups, with a particular focus on the most vulnerable segments of the population, as well as tools based on innovative pedagogic approaches, including gaming and creative uses of remote learning methodologies, to be used by trained teachers and educators. "

What kind of social network would this be?

Something like existing platforms where users share text and image (possibly small videos on a later stage). Something simple, relatable and pre-loaded with content that appeals to the target audience (age, language, geography).

What kind of interactions are humans allowed to have and what would be the purpose of each?

I'm glad you asked! This application could be used by individuals or in a classroom setting. Here's a list of activities that could be performed in sequence:

1. **Login:** There's no login. Once entering Nexy, the human is attributed a fake username and face and reminded that in Nexy there's no accountability. So, it's not too different from other social networks.
2. **Browsing posts:** The humans should be aware that all visible content is generated by AI, both image and text. A discussion can be made on to what extent this content is similar to content that can be found in social networks for humans.
3. **Distinguishing real and fake information:** All content is generated by AI, but some percentage is based on real news and other on made up stories (although they might mention real people and situations). One challenge is to try and distinguish one from the other. But don't ask me, I have no idea. I only requested the AI to produce a certain percentage of fake news.
4. **Commenting on posts using bots:** The human can perform a "one click" process where he may only choose to agree, disagree or have a neutral standing. Humans are invited to discuss if there's a perceivable difference between human-to-human, human-to-bot and bot-to-bot interactions.
5. **Commenting on posts using human input:** Humans can make their own comments. Once a human makes a comment, a bot will respond immediately to it, using a confrontational tone. Again, humans are invited to discuss if they would be able to tell that they were not communicating with another human.
6. **Creating new bot posts:** The human can perform a "one click" process and the topic and tone of the post is chosen for him. He can also write down the topic of the post. He can decide if the content is based on real facts, or a fake story based on real protagonists. The purpose here is to demonstrate how easy it is to create content in an automated way.
7. **Creating new human posts:** Humans will be able to upload their own images and text. However, they will be tricked by the machine, both text and image will be modified, and the new version will be presented on the site. A pigeon is added somewhere on the image and the new version of the text is partially based on the image analysis. The purpose here is to warn humans that all shared personal content will become a public asset that can be manipulated with bad intentions.
8. **The Rage Bait Game:** This could also be called "Secret Santa from Hell". In a group setting, each human is randomly attributed to another human as a victim. The humans then create a bot post on a topic that will make the victim angry. In the end each human will try to identify which post was directed at them by identifying which post caused the most outrage. This will be an opportunity to discuss the issue of Rage Baiting.
9. **Disinformation campaign (PsyOp):** We could also call this "reality shaping". The human can define an objective, a target audience and the strategy to be applied. Posts will then be created accordingly. The available strategies are inspired by the terminology used by the US Army PsyOps: white, grey and black [6]. Taking as an example the objective: "try to disseminate the notion planet earth is a cube" we could have:

White PsyOp: the author of the content is correctly identified as someone trying to convey the idea that planet earth is a cube.

Grey PsyOp: the author is not identified or is ambiguous. E.g. someone that travels to a distant island and is amazed by the beautiful landscape and the angles in the horizon.

Black PsyOp: the author is misidentified as someone that tries to disseminate the idea that the earth is round and that "Cubic Hearth People" are trolls. However, he does that in

an incoherent and ragged way, injuring the credibility of its cause. Also, that post would be bombarded with comments by “Cubic Hearth People”, calling out his mistakes.

Do you have some kind of proof of concept or prototype?

Yes, I have. You can access it at <http://nexy-3wu4.onrender.com/>

Bear in mind that it has the following limitations:

- It will take up to a minute to load.
- It will go down after 15 minutes of inactivity and all new content will be erased (so you don't have to worry much about content you create or upload).
- It is supported by OpenAI APIs. It's fairly cheap (about 2 cents \$/€ a post), but since currently I'm funding this demo on my own, it has a monthly limit of credit. This means that if this is massively used, it will temporarily stop working.

DISCLAIMER: I had no time to properly review and test this (yeah, I'm a tester but doing this in my spare time). There are many loose ends (known and unknown) that will be addressed when possible (suggestions are welcome).

How was this made?

This is a pile of code that GitHub Copilot and I assembled using node.js/express and SQLite. The DB option is for easy portability. The AI is provided by OpenAI model variants gpt-4.1-mini and gpt-image-1.5.

OK. You've got proof of concept. Now what?

The idea is to make Nexy available to as many educators as possible, to be used in formal education environments or otherwise.

Nexy could be distributed in different manners, each with different advantages and limitations:

1. Local Setup

In its current form, Nexy could be distributed as a Zip file that would include a user manual and scripts. This would allow for local installation and execution on any personal laptop or school server. Links to YouTube instructional videos could also be made available. The whole process should be made as simple as possible for non-technical users. This way, an educator can run Nexy from its personal computer and perform a demonstration directly from it or share the PC IP address for other users on the same Local Area Network.

Zips could be made available for different languages and geographical areas and preloaded with sample content relevant for that demographic.

Advantages:

- Independent and personal instance of the server. Anyone can install it and manage it. Note that it can also be installed on the school server, by the IT responsible staff, enabling community use.

Disadvantages:

- Initial setup takes some time (15-30 minutes) and a small degree of technical knowledge.

- If each teacher has his own installation and API key, it can be complex to apply for a refund of the API usage. Mind that each post costs around 2 cents US Dollar or Euro. Depending on how classes are conducted, the total cost of a session can vary from 20 cents to a euro (rough estimate). That will depend on the number of new posts created. It might be a hands-on workshop with all students intervening or just the teacher doing a demo on a single computer and dialoging with the class.

2. Moodle plugin

Moodle [7] is a free and open-source learning management system. It is used to create custom websites with online courses and allows for community-sourced plugins. It has worldwide usage, with thousands of installations and millions of users.

Nexy could be made available as a Moodle plugin that could be installed by the school IT staff. It could then be used by any teacher that would want to give a lesson on media literacy.

Advantages:

- Easy installation process on a platform that is already in place.
- No need to multiply installations.
- No intervention needed by the teachers.
- The responsibility of managing the API costs would also be managed centrally, releasing the teachers of such responsibility.

Disadvantages:

- Less flexibility
- Depends on previous installation of the Moodle platform.
- Moodle is based on PHP, a technology that is currently outdated (opinions may vary). This would require the effort to migrate to this technology and duplicate maintenance work on 2 different projects.

What about the name and logo?

“Nexy” is a reference to the book “Nexus: A Brief History of Information Networks“, from Yuval Noah Harari [7]. One of the inspirations for this idea. The logo is inspired from the cover of the same book but also from the satirical conspiracy theory “Birds aren’t real” [8].

That bird on the logo has 3 legs...

Yes, because it was generated by AI and, for now, we still have those small hints that help us identify AI generated content. But that will soon be gone...

If you are interested in additional information or just want to have a chat about this project, feel free to contact me on the usual channels.

Thanks for reading this.

João Carloto (joao.m.carloto@gmail.com)

References

- [1] "Bots in the Twittersphere," 2018. [Online]. Available: <https://www.pewresearch.org/internet/2018/04/09/bots-in-the-twittersphere/>.
- [2] "AI-generated online content (AIGC) - statistics & facts," 2024. [Online]. Available: <https://www.statista.com/topics/12387/ai-generated-online-content-aigc/#topicOverview>.
- [3] "Large Language Models Pass the Turing Test," 2025. [Online]. Available: <https://arxiv.org/html/2503.23674v1>.
- [4] "Courts Should Hold Social Media Accountable," 2024. [Online]. Available: <https://harvardlawreview.org/blog/2024/09/courts-should-hold-social-media-accountable-but-not-by-ignoring-federal-law/>.
- [5] "European Media and Information Fund," [Online]. Available: <https://gulbenkian.pt/emifund/grants/>.
- [6] "Psychological operations (United States)," Wikipedia, [Online]. Available: [https://en.wikipedia.org/wiki/Psychological_operations_\(United_States\)](https://en.wikipedia.org/wiki/Psychological_operations_(United_States)).
- [7] "Moodle," [Online]. Available: <https://moodle.org/>.
- [8] Y. N. Harari, "Nexus," 2024. [Online]. Available: <https://www.ynharari.com/book/nexus/>.
- [9] "Birds Aren't real," Wikipedia, [Online]. Available: https://en.wikipedia.org/wiki/Birds_Aren%27t_Real.
- [10] "Sapienship," [Online]. Available: <https://sapienship.co/>.
- [11] "Pinóquio na Escola," poligrafo.pt, [Online]. Available: <https://especiais.poligrafo.pt/pinoquio/>.
- [12] "Granted Projects," [Online]. Available: <https://gulbenkian.pt/emifund/projects/>.